



浙江大学爱丁堡大学联合学院

ZJU-UoE Institute

## Lecture 2 - Geometric image transformations

---

Nicola Romanò - [nicola.romano@ed.ac.uk](mailto:nicola.romano@ed.ac.uk)

## **A brief recap**

---

Last time we learnt how to open and display images using Python.



General plotting library



Specific functions for image  
manipulation



Library for matrix operations





```
import matplotlib.pyplot as plt
img = plt.imread("cells.jpg")
plt.imshow(img)
plt.show()
```



```
from skimage import io
img = io.imread("cells.jpg")
io.imshow(img)
io.show()
```

## Colour mapping

```
fig, ax = plt.subplots(ncols=2, nrows=2)
ax[0, 0].imshow(img, cmap="gray")
ax[0, 1].imshow(img, cmap="viridis")
ax[1, 0].imshow(img, cmap="Greens")
ax[1, 1].imshow(img, cmap="PuBu")
plt.show()
```



Check out the Matplotlib website for a list of colourmaps!

<https://matplotlib.org/stable/tutorials/colors/colormaps.html>

- Describe use cases for geometric transformations of images
- Use Python to crop images (in 2 or more dimensions)
- Explain the theory of affine transformations
- Implement basic transformations in Python (translation, scaling, rotation)



## **Geometric image transformations**

---

**Geometric image transformations** are operations that change the image geometry without altering its pixel values (mostly...).

Examples are cropping, translating, scaling rotating an image.




**Geometric image transformations** are operations that change the image geometry without altering its pixel values (mostly...).

Examples are cropping, translating, scaling rotating an image.

Example use cases:

- Analysing only part of an image (cropping)
- Making sure all input images for a pipeline are the same size (scaling, cropping)
- Aligning multiple images (e.g. in video stabilization) (rotating, translating)
- ...




## **Geometric image transformations**

---

### **Cropping**

Cropping is as easy as subsetting the image matrix.

```
# img contains a 512 x 512 image  
# Take the top-left 100 x 100 pixels  
img1 = img[0:100, 0:100] # shape (100, 100)
```




`img[0:100, 0:100]`

Cropping is as easy as subsetting the image matrix.

```
# img contains a 512 x 512 image
# Take the top-left 100 x 100 pixels
img1 = img[0:100, 0:100] # shape (100, 100)

# A 100 pixel wide strip on the left...
img2 = img[:, 0:100] # shape (512, 100)
# ...or on the right!
img3 = img[:, -100:] # shape (512, 100)
```




`img[:, 0:100]`

`img[:, -100:]`

## Cropping in more than two dimensions

Since images are just tensors, we can crop them in any dimension.

```
# img is a 512 x 512 RGB image
# img.shape is (3, 512, 512)
# Extract the green channel
img_green = img[1] # Shape (512, 512)
```




## Cropping in more than two dimensions

Since images are just tensors, we can crop them in any dimension.

```
# img is a 512 x 512 RGB image
# img.shape is (3, 512, 512)
# Extract the green channel
img_green = img[1] # Shape (512, 512)

# video is a 512 x 512 video of 300 frames
# video.shape is (300, 512, 512)
# Take the first 50 frames
video_crop = video[0:50]
```



We have a 100 frames video of a 512 x 512 z-stack with 60 planes loaded in `zstack`.

`zstack.shape` is (100, 60, 512, 512)


**What does this code give you?**

```
result = stack[50:70, :, 100:300]
```



## Affine transformations

In Euclidean geometry, an affine transformation is a geometric transformation that preserves lines and parallelism but not necessarily distances and angles. [Wikipedia]



We want to transform  $P(x; y)$  into  $P'(x'; y')$ .


$$P' : \begin{cases} x' = f(x, y) \\ y' = g(x, y) \end{cases}$$

Where  $f$  and  $g$  are linear functions.



# Affine transformations

In Euclidean geometry, an affine transformation is a geometric transformation that preserves lines and parallelism but not necessarily distances and angles. [Wikipedia]




We want to transform  $P(x; y)$  into  $P'(x'; y')$ .

$$P' : \begin{cases} x' = f(x, y) \\ y' = g(x, y) \end{cases}$$

Where  $f$  and  $g$  are linear functions.

We can generalise this to an image, by applying the transformation to every pixel.




## **Geometric image transformations**


---

### **Scaling**

Scaling transforms the coordinates as:

$$\begin{cases} x' = s_x * x \\ y' = s_y * y \end{cases}$$





Scaling transforms the coordinates as:


$$\begin{cases} x' = s_x * x \\ y' = s_y * y \end{cases}$$

We can write this in matrix form\*:

$$\begin{bmatrix} x' \\ y' \end{bmatrix} = \begin{bmatrix} s_x & 0 \\ 0 & s_y \end{bmatrix} \begin{bmatrix} x \\ y \end{bmatrix}$$

This is called the scaling **transformation matrix**.

\* Don't remember how matrix multiplication works? Check out [Wikipedia!](#)



Two problems:

When **upscaling** we need to generate new information.


When **downscaling** we need to decide "how to lose" information.

**Interpolation** is the solution!



## Nearest-neighbour interpolation

The simplest way to resize an image is to use nearest-neighbour interpolation. Each pixel of the scaled image will have the colour of the nearest pixel in the original image.

In this "1D" example, we resize a 1x2 image to 1x6





## Nearest-neighbour interpolation



Upscaling of a 5x5 image by a factor of 2, to get a 10x10 image, with nearest neighbour interpolation

## Nearest-neighbour interpolation



Upscaling of a 5x5 image by a factor of 2, to get a 10x10 image, with nearest neighbour interpolation



```
from skimage.transform import rescale  
img_scaled = rescale(img, 2, order=0)
```



## Scaling with interpolation


Better ways to scale an image involve changing the pixel values of the rescaled image based on their neighbourhood.

For example we could use linear interpolation



## Linear interpolation


The same applies in 2D, although we need to take into account the values of both horizontal and vertical neighbours (**bilinear interpolation**).



Upscaling of a 5x5 image by a factor of 2, to get a 10x10 image, with bi-linear interpolation

```
from skimage.transform import rescale  
img_scaled = rescale(img, 2, order=1)
```


## Nearest neighbour vs linear interpolation



Comparison of nearest neighbour and linear interpolation - Source: Wikipedia

# Higher orders of interpolation

We can use higher orders of interpolation to produce smoother results.



Comparison of nearest neighbour and linear interpolation - Source: Wikipedia

Scikit Image supports values from 0 to 5 in the `order` parameter of the `rescale` function. 0 is nearest neighbour, 1 is bi-linear, 2 is bi-quadratic and so on.


To scale to a target size, rather than by a specific factor, we can use `resize` instead of `rescale`.

```
import matplotlib.pyplot as plt
from skimage.transform import resize, rescale

img = plt.imread("cells.jpg")
img_scaled1 = rescale(img, 2)
img_scaled2 = resize(img, (150, 150))
```


## Local mean downscaling

Another simple solution for downscaling is calculating the local mean of each pixel



## Local mean downscaling

Another simple solution for downscaling is calculating the local mean of each pixel



```
from skimage.transform import downscale_local_mean  
img_small = downscale_local_mean(img, (2,2))
```

Image needs to be padded if the size is not a multiple of the downscaling factor.  
This method is fast but not good at keeping fine details.

$$\begin{bmatrix} x' \\ y' \end{bmatrix} = \begin{bmatrix} s_x & 0 \\ 0 & s_y \end{bmatrix} \begin{bmatrix} x \\ y \end{bmatrix}$$

`skimage.transform.rescale` → scales an image by a specific factor (>1 upscaling; <1 downscaling.). Can specify a different scaling factor for each dimension of the image.

`skimage.transform.resize` → scales an image to a target size.

`skimage.transform.downscale_local_mean` → downscales the image by a specific factor (>1), using the local mean of each pixel.




## **Geometric image transformations**

---


### **Rotations**

## Rotating a point around the origin

We want to rotate a point  $P(x, y)$  around the origin by an angle  $\theta$  to get  $P'(x', y')$ .




## Rotating a point around the origin



$$P : \begin{cases} x = r \cos(\alpha) \\ y = r \sin(\alpha) \end{cases} \quad P' : \begin{cases} x' = r \cos(\alpha + \theta) \\ y' = r \sin(\alpha + \theta) \end{cases}$$


## Rotating a point around the origin



$$P : \begin{cases} x = r \cos(\alpha) \\ y = r \sin(\alpha) \end{cases} \quad P' : \begin{cases} x' = r \cos(\alpha + \theta) = r \cos(\alpha)\cos(\theta) - r \sin(\alpha)\sin(\theta) \\ y' = r \sin(\alpha + \theta) = r \sin(\alpha)\cos(\theta) + r \sin(\theta)\cos(\alpha) \end{cases}$$

Why? [Wikipedia to the rescue!](#)

## Rotating a point around the origin




$$P : \begin{cases} x = r \cos(\alpha) \\ y = r \sin(\alpha) \end{cases} \quad P' : \begin{cases} x' = r \cos(\alpha + \theta) = r \cos(\alpha)\cos(\theta) - r \sin(\alpha)\sin(\theta) \\ y' = r \sin(\alpha + \theta) = r \sin(\alpha)\cos(\theta) + r \sin(\theta)\cos(\alpha) \end{cases}$$

Why? Wikipedia to the rescue!

Thus:

$$P' : \begin{cases} x' = x \cos(\theta) - y \sin(\theta) \\ y' = y \cos(\theta) + x \sin(\theta) \end{cases}$$

## Rotating a point around the origin



$$P : \begin{cases} x = r \cos(\alpha) \\ y = r \sin(\alpha) \end{cases} \quad P' : \begin{cases} x' = r \cos(\alpha + \theta) = r \cos(\alpha)\cos(\theta) - r \sin(\alpha)\sin(\theta) \\ y' = r \sin(\alpha + \theta) = r \sin(\alpha)\cos(\theta) + r \sin(\theta)\cos(\alpha) \end{cases}$$

Why? Wikipedia to the rescue!

Thus:

$$P' : \begin{cases} x' = x \cos(\theta) - y \sin(\theta) \\ y' = y \cos(\theta) + x \sin(\theta) \end{cases}$$

$$\begin{bmatrix} x' \\ y' \end{bmatrix} = \begin{bmatrix} \cos\theta & -\sin\theta \\ \sin\theta & \cos\theta \end{bmatrix} \begin{bmatrix} x \\ y \end{bmatrix}$$

Rotation transformation matrix

## Rotating in Scikit Image

To rotate an image we will:

- Offset our image so that it is centered on the origin
- Generate a rotation matrix
- Multiply the coordinates of each pixel by the rotation matrix
- Shift back the image to its original position

## Rotating in Scikit Image

To rotate an image we will:



- Offset our image so that it is centered on the origin
- Generate a rotation matrix
- Multiply the coordinates of each pixel by the rotation matrix
- Shift back the image to its original position

Luckily, Scikit Image has a function for that, `skimage.transform.rotate`.

```
import matplotlib.pyplot as plt
from skimage.transform import rotate

img = plt.imread("cells.jpg")
img_rotated = rotate(img, 20)

plt.imshow(img_rotated, cmap="gray")
plt.show()
```



Note: we lost part of the image and we "gained" black pixels around it.



## **Geometric image transformations**

---

### **Translation**

## Translation

To translate we offset the points by  $(t_x, t_y)$ :  $\begin{cases} x' = x + t_x \\ y' = y + t_y \end{cases}$

The transformation matrix is different from the ones we saw for before:

$$\begin{bmatrix} x' \\ y' \end{bmatrix} = \begin{bmatrix} x \\ y \end{bmatrix} + \begin{bmatrix} t_x \\ t_y \end{bmatrix}$$

## Translation

To translate we offset the points by  $(t_x, t_y)$ : 
$$\begin{cases} x' = x + t_x \\ y' = y + t_y \end{cases}$$

The transformation matrix is different from the ones we saw for before:

$$\begin{bmatrix} x' \\ y' \end{bmatrix} = \begin{bmatrix} x \\ y \end{bmatrix} + \begin{bmatrix} t_x \\ t_y \end{bmatrix}$$

We can generate a transformation matrix using `SimilarityTransform` and apply it using `warp`.

```
from skimage.transform import SimilarityTransform, warp
# Create a 10x10 white image
img = np.ones(shape=(10, 10))
# Make the central pixels black
img[5:7, 5:7] = 0

m = SimilarityTransform(translation = (2, 5))
img_translated = warp(img, m)
```

# Translation

To translate we offset the points by  $(t_x, t_y)$ : 
$$\begin{cases} x' = x + t_x \\ y' = y + t_y \end{cases}$$


The transformation matrix is different from the ones we saw for before:

$$\begin{bmatrix} x' \\ y' \end{bmatrix} = \begin{bmatrix} x \\ y \end{bmatrix} + \begin{bmatrix} t_x \\ t_y \end{bmatrix}$$

We can generate a transformation matrix using `SimilarityTransform` and apply it using `warp`.

```
from skimage.transform import SimilarityTransform, warp
# Create a 10x10 white image
img = np.ones(shape=(10, 10))
# Make the central pixels black
img[5:7, 5:7] = 0

m = SimilarityTransform(translation = (2, 5))
img_translated = warp(img, m)
```



### Translation

$$\begin{bmatrix} x' \\ y' \end{bmatrix} = \begin{bmatrix} x \\ y \end{bmatrix} + \begin{bmatrix} t_x \\ t_y \end{bmatrix}$$

### Rotation

$$\begin{bmatrix} x' \\ y' \end{bmatrix} = \begin{bmatrix} \cos\theta & -\sin\theta \\ \sin\theta & \cos\theta \end{bmatrix} \begin{bmatrix} x \\ y \end{bmatrix}$$

### Scaling

$$\begin{bmatrix} x' \\ y' \end{bmatrix} = \begin{bmatrix} s_x & 0 \\ 0 & s_y \end{bmatrix} \begin{bmatrix} x \\ y \end{bmatrix}$$

## The affine transformation matrix

Sometimes we want to combine rotation, scaling and translation into a single operation. We can rewrite the matrices as such (these are called homogeneous coordinates):

### Translation

$$\begin{bmatrix} x' \\ y' \\ 1 \end{bmatrix} = \begin{bmatrix} 1 & 0 & t_x \\ 0 & 1 & t_y \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} x \\ y \\ 1 \end{bmatrix}$$

### Rotation

$$\begin{bmatrix} x' \\ y' \\ 1 \end{bmatrix} = \begin{bmatrix} \cos\theta & -\sin\theta & 0 \\ \sin\theta & \cos\theta & 0 \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} x \\ y \\ 1 \end{bmatrix}$$

### Scaling

$$\begin{bmatrix} x' \\ y' \\ 1 \end{bmatrix} = \begin{bmatrix} s_x & 0 & 0 \\ 0 & s_y & 0 \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} x \\ y \\ 1 \end{bmatrix}$$

We can now easily combine these 3x3 matrices by multiplying them together.

Remember that the order is important!

$A * B$  is **not** the same as  $B * A$ !

You can use `skimage.transform.SimilarityTransform` to create transformation matrices and combine them using the `+` operator (Note that under the hood this performs matrix multiplication!).

Using `skimage.transform.warp` we can apply these transformation matrices to an image.

## Optional exercise - want to try by yourself?

It might be interesting to try to code image rotation by yourself.

Hints:

- Generate the transformation matrices using `skimage.transform.SimilarityTransform`
- You will need three matrices: one translation matrix to offset your image by `(-xcenter, -ycenter)`; a rotation matrix to rotate your image around `(0, 0)`; and finally another translation matrix to translate your image back to its original position.
- You can combine the matrices as `m1+m2+m3`
- You can use `skimage.transform.warp` to apply the combined transformation matrix to your image.

*If you are stuck, try to look at the [source code for rotate!](#)*



- Affine transformations are simple yet powerful way to modify an image.
- Scikit Image allows generation of any transformation matrix...
- ...but also provides several pre-made functions for rotating, scaling, etc
- **Workshop 1** will allow you to practice what learned so far!