




浙江大学爱丁堡大学联合学院  
ZJU-UoE Institute

## Lecture 4 - Filters

---

Nicola Romanò - [nicola.romano@ed.ac.uk](mailto:nicola.romano@ed.ac.uk)

- Define rank and convolutional filters
- Explain their use in image analysis
- Implement basic filters in Python



Operations for manipulating pixel intensities

Two types of operations:

- Point operations - Change pixel intensity based only on its value -  $I'_{(x,y)} = f(I_{x,y})$  (see Lecture 3)
- **Neighbourhood operations** - Change pixel intensity based on the intensity of the pixel and its neighbours.

Neighbourhood operations, often called **filters** allow to modify an image in a way that is not possible with point operations.

- Detect simple structures such as edges, corners, lines, etc.
- Perform operations such as smoothing, sharpening, etc.
- Noise reduction

Neighbourhood operations, often called **filters** allow to modify an image in a way that is not possible with point operations.


- Detect simple structures such as edges, corners, lines, etc.
- Perform operations such as smoothing, sharpening, etc.
- Noise reduction

Today we will look at:

- **Rank filters** - the new pixel value is a function of the rank of the pixel values of the neighbourhood
- **Convolutional filters** - the new pixel value is a weighted sum of the pixel values of the neighbourhood



## Rank filters

---



- We decide on a window size  $w \times h$  (other non-rectangular shapes are possible)
- We traverse each pixel in the image and take its  $w \times h$  neighbourhood
- We rank the intensity of each pixel in the neighbourhood
- We take a specific value (e.g. minimum, maximum, median) and set the output value to this value

## Example - median filter





## Rank filters in Scikit Image

Rank filters are implemented in Scikit Image in the 'skimage.rank' module.

These filters require a *footprint* of the pixel neighbourhood, as a matrix of 0 and 1.


For example

```
from skimage.rank import median  
  
# 3x3 neighbourhood  
img_median = median(img, footprint = np.ones(3, 3))
```


`np.ones(3, 3)` creates

$$\begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix}$$


Original




Footprint: 3x3



Footprint: 7x7



Footprint: 3x15



## Rank filters in Scikit Image

Rank filters are implemented in Scikit Image in the 'skimage.rank' module.

These filters require a *footprint* of the pixel neighbourhood, as a matrix of 0 and 1.


For example

```
from skimage.rank import median  
  
# 3x3 neighbourhood  
img_median = median(img, footprint = np.ones(3, 3))
```


`np.ones(3, 3)` creates

$$\begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix}$$


Original




Footprint: 3x3



Footprint: 7x7



Footprint: 3x15



What happens with larger footprints? What about non-square footprints?


## Non-rectangular footprints

The footprint can be a *non-rectangular* matrix. For example, you can generate a **circular** shape using the `skimage.morphology.disk` function or a **diamond** shape using the `skimage.morphology.diamond` function.

```
from skimage.morphology import disk, diamond

dsk = disk(5) # A disk, radius 5 px
dia = diamond(5) # A diamond, radius 5 px

img_disk = median(img, footprint=dsk)
img_diamond = median(img, footprint=dia)
```




The difference in output can be subtle, but it can be useful in some cases to use a different shape.

Rank filters are very simple, but have useful applications.  
The **median filter** is used to remove noise and to smooth images.

Rank filters are very simple, but have useful applications.

The **median filter** is used to remove noise and to smooth images.

The **maximum filter** can be used in binary images to remove small "holes" (can you think of why?). It is also a very common filter used in modern neural networks for image analysis.




## Rank filters - use cases

Rank filters are very simple, but have useful applications.

The **median filter** is used to remove noise and to smooth images.

The **maximum filter** can be used in binary images to remove small "holes" (can you think of why?). It is also a very common filter used in modern neural networks for image analysis.

The **minimum filter** can be used to remove small bright spots (can you think of why?) .



## Convolutional filters

---

A **convolutional filter** consists of a small matrix, called a **kernel**, that is used to process an image. Convolution takes each pixel of the image together with its neighbours and adds them together, weighting each neighbour by the value of a kernel of the same size of the neighbourhood.



# Image

255	255	80	255	255
255	50	80	50	255
80	80	0	80	80
255	50	80	50	255
255	255	80	255	255

# Kernel

$$\begin{bmatrix} 0 & -1 & 0 \\ -1 & 5 & -1 \\ 0 & -1 & 0 \end{bmatrix}$$

# Image

255	255	80	255	255
255	50	80	50	255
80	80	0	80	80
255	50	80	50	255
255	255	80	255	255


# Kernel

$$\begin{bmatrix} 0 & -1 & 0 \\ -1 & 5 & -1 \\ 0 & -1 & 0 \end{bmatrix}$$

The convolved pixel value will be


$$255 * 0 + 50 * (-1) + 80 * 0 + 80 * (-1) + 80 * 5 + 0 * (-1) + 255 * 0 + 50 * (-1) + 80 * 0 = \mathbf{220}$$

## Example of a convolutional filter - result



Our image, after applying the convolutional filter.

## Example of a convolutional filter - result



Our image, after applying the convolutional filter.


Convolutional filters are **shift invariant**, meaning the result of the convolution is the same regardless of the position of the pixel in the original image (as long as it has the same neighbourhood).

## Common convolutional filters - averaging filter

The **averaging filter** or **box blur** is a simple filter that is used to reduce noise in an image. It simply takes the average of the pixel values in the neighbourhood.

Example 3x3 kernel:

$$\begin{bmatrix} 1/9 & 1/9 & 1/9 \\ 1/9 & 1/9 & 1/9 \\ 1/9 & 1/9 & 1/9 \end{bmatrix}$$





Not as good as the **median filter**, as it is sensitive to outliers and does not preserve edges as well (can you get an intuition as to why?).

## Common convolutional filters - Gaussian filter

The **Gaussian filter** is a filter that is used to smooth images.

It uses a **Gaussian** function to weight the pixel values in the neighbourhood.



$$G_{\sigma} = \frac{1}{2\pi\sigma^2} e^{-\frac{(x^2+y^2)}{2\sigma^2}}$$


Example of gaussian kernels 3x3 and 5x5 (approx.):

$$\frac{1}{16} \begin{bmatrix} 1 & 2 & 1 \\ 2 & 4 & 2 \\ 1 & 2 & 1 \end{bmatrix}$$

$$\frac{1}{256} \begin{bmatrix} 1 & 4 & 6 & 4 & 1 \\ 4 & 16 & 24 & 16 & 4 \\ 6 & 24 & 36 & 24 & 6 \\ 4 & 16 & 24 & 16 & 4 \\ 1 & 4 & 6 & 4 & 1 \end{bmatrix}$$

## Common convolutional filters - Gaussian filter - result

Example of gaussian filters with increasing sigma values:




Gaussian filters result in a blurring effect, which can be useful for removing noise, or in general removing the finer detail of the image (low-pass filtering).

```
from skimage.filters import gaussian  
img_blurred = gaussian(image, sigma=1)
```

## Sharpening filters

Sharpening filters are used to increase the detail of an image.




If we apply a Gaussian filter to the image we will remove detail, so if we subtract the result from the original image we will get the "details" of the image.



## Sharpening filters

Sharpening filters are used to increase the detail of an image.




We can now add the details to the original image, thus sharpening it!

This process is called "**unsharp masking**".

# Unsharp masking in Scikit Image

You can use the `skimage.filters.unsharp_mask` function to apply unsharp masking to images.

```
from skimage.filters import unsharp_mask  
img_sharpened = unsharp_mask(img, radius=15, amount=2)
```




Exploring the parameter space of unsharp masking.

# Unsharp masking in Scikit Image

You can use the `skimage.filters.unsharp_mask` function to apply unsharp masking to images.

```
from skimage.filters import unsharp_mask  
img_sharpened = unsharp_mask(img, radius=15, amount=2)
```

**Exercise:** try writing your own unsharp masking function.  
It should accept an image, a radius (the  $\sigma$  of the gaussian blur)  
and an amount for sharpening and return the sharpened image.



Exploring the parameter space of unsharp masking.

## What happens at the edges?

Convolutional filters traverse the image pixel by pixel and apply a function that takes into account the pixel's neighbourhood.

**What happens at the edges**, where the neighbourhood is not complete?  
There are several ways to deal with this.


## What happens at the edges?

Convolutional filters traverse the image pixel by pixel and apply a function that takes into account the pixel's neighbourhood.

**What happens at the edges**, where the neighbourhood is not complete?

There are several ways to deal with this.

- **Zero-pad the image:** we create an extra padding of pixels with a value of zero around the image.



Zero-pad


## What happens at the edges?

Convolutional filters traverse the image pixel by pixel and apply a function that takes into account the pixel's neighbourhood.


**What happens at the edges**, where the neighbourhood is not complete?

There are several ways to deal with this.

- **Zero-pad the image:** we create an extra padding of pixels with a value of zero around the image.
- **Extend the image:** edge pixels are copied, corner pixels are repeated as *wedges*.



Zero-pad



Extend


## What happens at the edges?

Convolutional filters traverse the image pixel by pixel and apply a function that takes into account the pixel's neighbourhood.


**What happens at the edges**, where the neighbourhood is not complete?

There are several ways to deal with this.


- **Zero-pad the image:** we create an extra padding of pixels with a value of zero around the image.
- **Extend the image:** edge pixels are copied, corner pixels are repeated as *wedges*.
- **Wrap the image:** extra pixels are taken from the opposite side of the image.



Zero-pad



Extend



Wrap


## What happens at the edges?

Convolutional filters traverse the image pixel by pixel and apply a function that takes into account the pixel's neighbourhood.


**What happens at the edges**, where the neighbourhood is not complete?

There are several ways to deal with this.


- **Zero-pad the image:** we create an extra padding of pixels with a value of zero around the image.
- **Extend the image:** edge pixels are copied, corner pixels are repeated as *wedges*.
- **Wrap the image:** extra pixels are taken from the opposite side of the image.
- **Mirror the image:** extra pixels are created by mirroring the image.




Zero-pad



Extend



Wrap



Mirror




# What happens at the edges?

Convolutional filters traverse the image pixel by pixel and apply a function that takes into account the pixel's neighbourhood.


**What happens at the edges**, where the neighbourhood is not complete?

There are several ways to deal with this.


- **Zero-pad the image:** we create an extra padding of pixels with a value of zero around the image.
- **Extend the image:** edge pixels are copied, corner pixels are repeated as *wedges*.
- **Wrap the image:** extra pixels are taken from the opposite side of the image.
- **Mirror the image:** extra pixels are created by mirroring the image.
- **Crop the image:** we ignore the edge pixels. This will result in a smaller output image.
- **Crop the kernel:** we only use the kernel values corresponding to a pixel in the image. Kernel weights are adjusted to account for this.




Zero-pad




Extend




Wrap



Mirror



Crop the image



Crop the kernel

## Edge behaviour in Scikit Image

Filters in Scikit Image allow you to choose what to do with the edge pixels using the `mode` parameter.

For example

```
from skimage.filters import gaussian
image_smoothed = gaussian(image, sigma=1, mode="mirror")
```

Allowed values for `mode` are

- `constant` - padding with the value specified by `cval`
- `nearest` - extend
- `mirror` and `reflect` - both of these mirror the image, but `reflect` does not duplicate the edge pixel
- `wrap`

## Want to try out more?

You can design your own convolutional filter and apply it to an image using the `skimage.filters.edges.convolve` function.

What happens using this kernel?

$$\begin{bmatrix} 0 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 0 \end{bmatrix}$$

What about this?

$$\begin{bmatrix} 0 & 0 & 0 \\ 0 & 1.5 & 0 \\ 0 & 0 & 0 \end{bmatrix}$$

You can see convolutional filters in action at [this website](#).

- Image filters are a simple yet powerful way of manipulating images.
- Filters can be used to remove noise, smooth images, sharpen them, and more.
- Rank filters such as median, maximum and minimum work by choosing a specific value of the pixels in the neighbourhood, depending on their rank.
- Convolutional filters use a kernel to apply a linear transformation to the image.
- In the next lecture we will look at using filters for edge detection.