



浙江大学爱丁堡大学联合学院

ZJU-UoE Institute

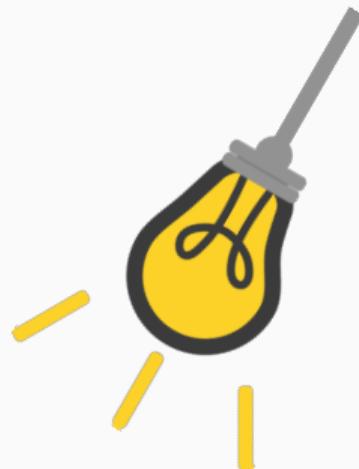
## Lecture 3 - Image histograms

---

Nicola Romanò - [nicola.romano@ed.ac.uk](mailto:nicola.romano@ed.ac.uk)

## Learning objectives

- Define and produce an image histogram
- Use histograms to interpret image quality
- Apply point operations to manipulate image histograms



## Histograms

---

## What is a histogram?

"A histogram is an approximate representation of the distribution of numerical data." ([Wikipedia](#))

Given a variable  $x$ , in the interval  $[x_{min}; x_{max}]$ , we divide this interval (or a suitably larger one) into non-overlapping **bins** and count the number of values of  $x$  falling into each bin.

Generally, we choose bins to be of equal size.

## Histogram example

Given  $x = 5, 7, 10, 22, 35, 88, 26, 74, 22, 95$

We can choose the interval  $[0; 99]$ , and divide it into bins of size 10  $[0 : 9][10; 19] \dots [90 : 99]$

## Histogram example

Given  $x = 5, 7, 10, 22, 35, 88, 26, 74, 22, 95$

We can choose the interval  $[0; 99]$ , and divide it into bins of size 10  $[0 : 9][10; 19] \dots [90 : 99]$

We now count the occurrences of  $x$  in each bin.

$$[0 : 9] \rightarrow 2$$

$$[10 : 19] \rightarrow 1$$

...and so on

## Histogram example

Given  $x = 5, 7, 10, 22, 35, 88, 26, 74, 22, 95$

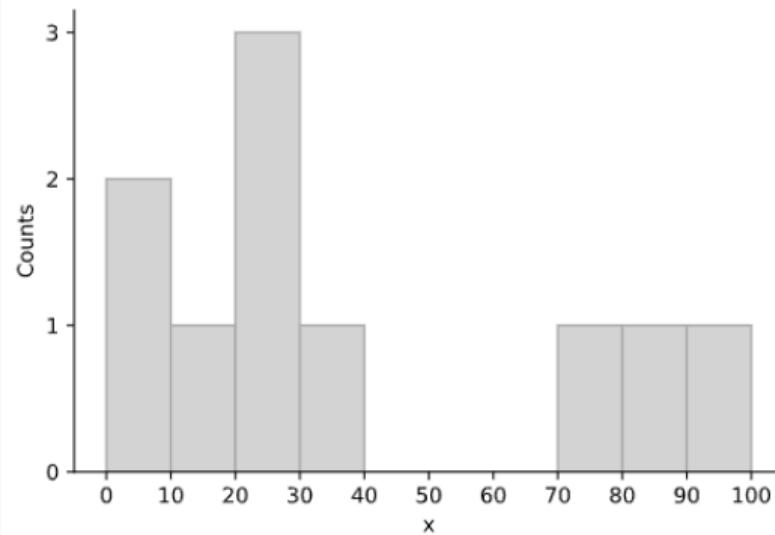
We can choose the interval  $[0 : 99]$ , and divide it into bins of size 10  $[0 : 9][10 : 19] \dots [90 : 99]$

We now count the occurrences of  $x$  in each bin.

$$[0 : 9] \rightarrow 2$$

$$[10 : 19] \rightarrow 1$$

...and so on



$x =$ 

|   |   |    |    |    |    |    |    |    |    |
|---|---|----|----|----|----|----|----|----|----|
| 5 | 7 | 10 | 22 | 35 | 88 | 26 | 74 | 22 | 95 |
|---|---|----|----|----|----|----|----|----|----|

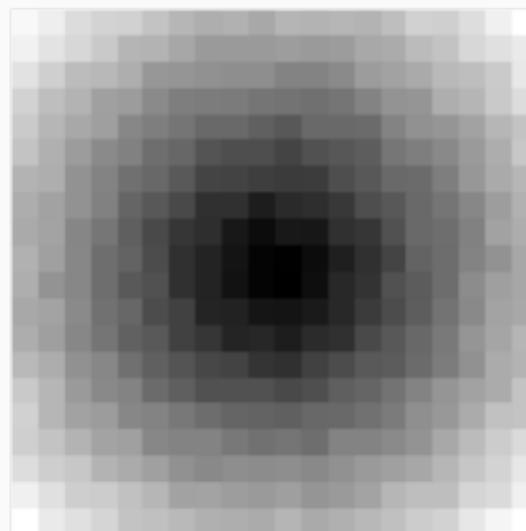
## Histogram of an image

Images are just matrices, so we can count the occurrences of each pixel value in the image.

### Example

An 8-bit image, 20 x 20

```
> print(img.shape)
(20, 20)
>print(np.min(img))
13
>print(np.max(img))
177
```



# Histogram of an image

Images are just matrices, so we can count the occurrences of each pixel value in the image.

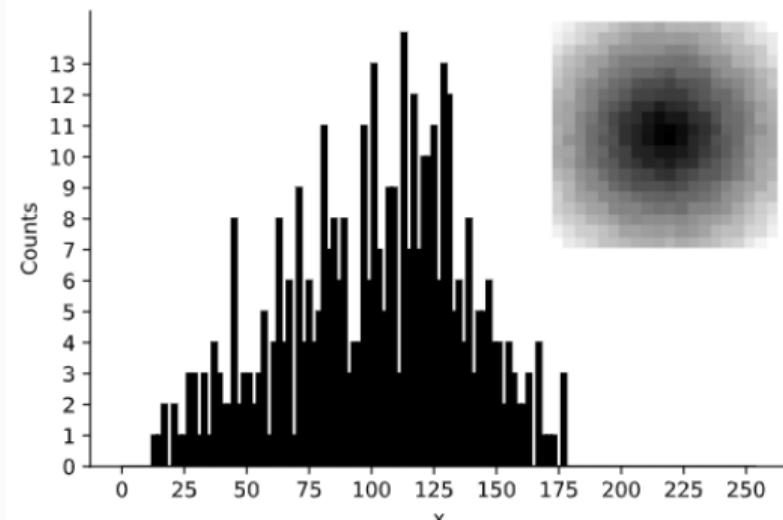
## Example

An 8-bit image, 20 x 20

```
> print(img.shape)
(20, 20)
> print(np.min(img))
13
> print(np.max(img))
177
```

Maximum dynamic range [0, 255]

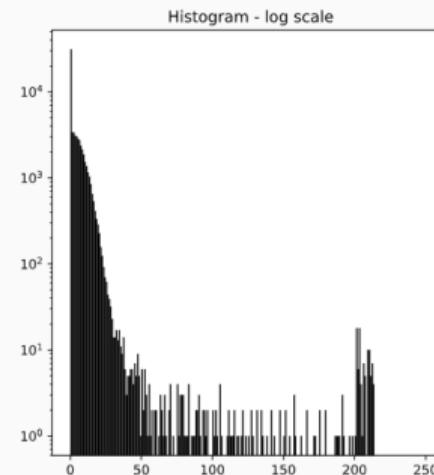
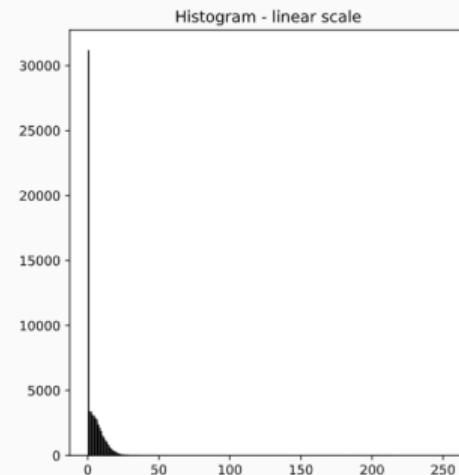
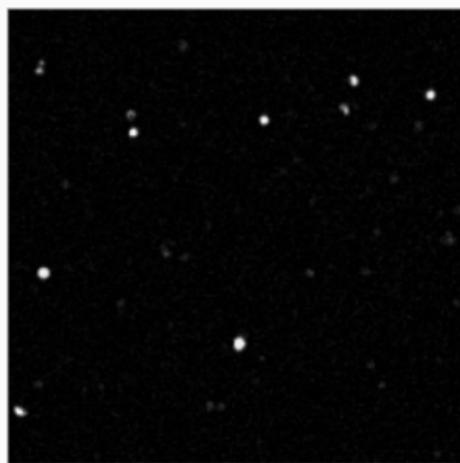
We count the occurrences of each pixel with intensity between 0 and 255.



What can you tell about the image from looking at the histogram?

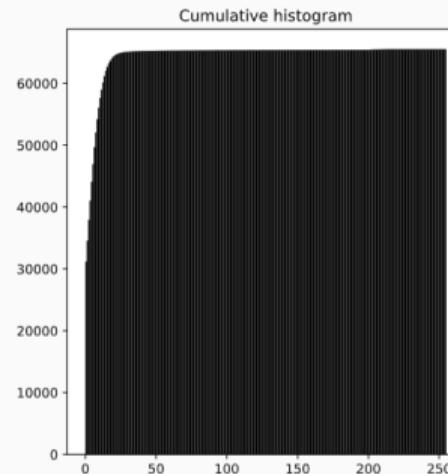
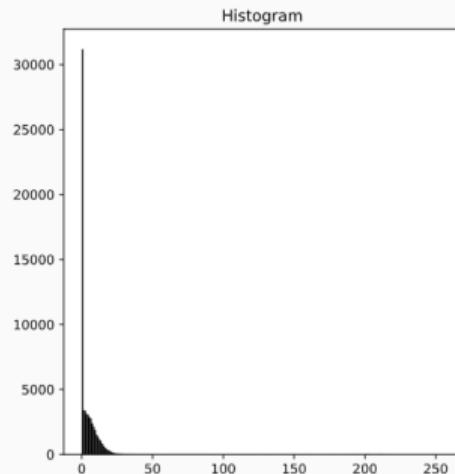
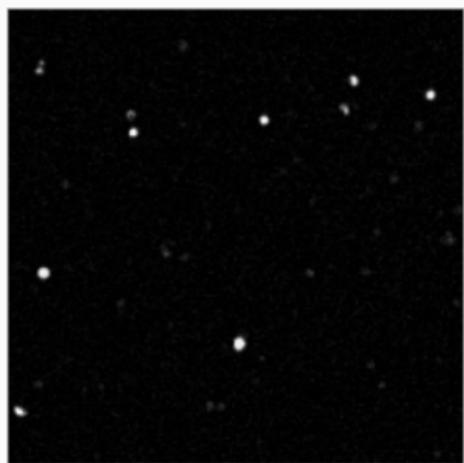
## Logarithmic scale

We can plot the histogram on a logarithmic scale. Useful when images contain many pixels of a specific intensity (e.g. cells on a black background).



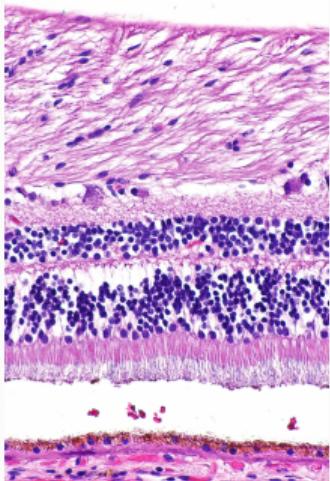
## Cumulative histogram

Alternatively, we can plot a cumulative histogram (see later for uses). Each bar is the count of pixels with intensity between 0 and the corresponding bin.

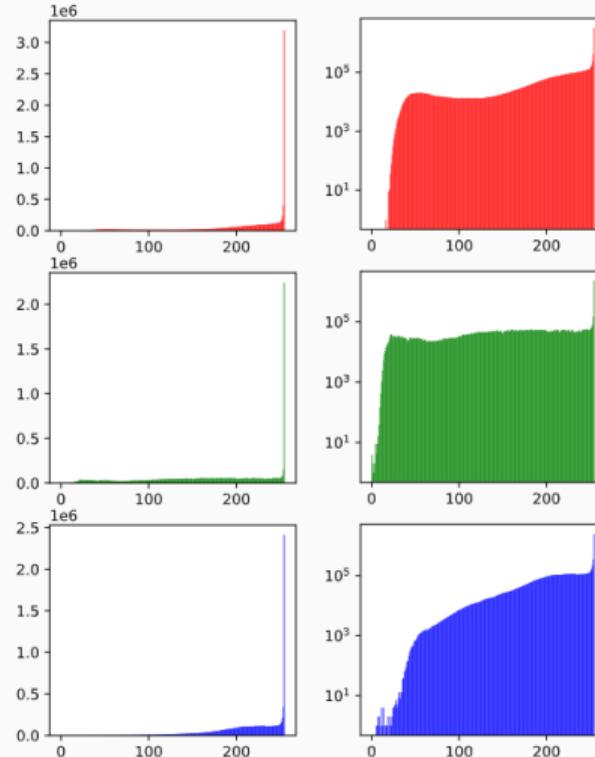


## Histograms of RGB images

For RGB (or multi-channel) images, it makes sense to have a single histogram per channel



H&E staining of retina (cell nuclei stained blue-purple  
and extracellular material stained pink)  
Librepath - CC BY-SA 3.0



RGB histogram - left linear, right logarithmic.

## Histograms are destructive

It is important to remember that producing a histogram is a many-to-one operation! We can generate a histogram from an image, but cannot recreate the image from the histogram.



These three images have the same histogram!

## Drawing a histogram using Matplotlib

The Matplotlib `hist` function makes it very easy to plot a histogram

```
import matplotlib.pyplot as plt
img = io.imread("cells.tif")

fig, ax = plt.subplots(1, 2, figsize=(12, 6))
# Linear
ax[0].hist(img.ravel(), bins=range(255), color="black")
ax[0].set_title("Histogram - linear scale")
```

The `bins` parameter accepts either a sequence defining the edges of the bins, or a single value with the number of desired bins.

## Drawing a histogram using Matplotlib

The Matplotlib `hist` function makes it very easy to plot a histogram

```
import matplotlib.pyplot as plt
img = io.imread("cells.tif")

fig, ax = plt.subplots(1, 2, figsize=(12, 6))
# Linear
ax[0].hist(img.ravel(), bins=range(255), color="black")
ax[0].set_title("Histogram - linear scale")
# Log
ax[1].hist(img.ravel(), bins=range(255), color="black", log=True)
ax[1].set_title("Histogram - log scale")
plt.show()
```

The `bins` parameter accepts either a sequence defining the edges of the bins, or a single value with the number of desired bins.

## **Brightness and contrast**

---

## Brightness

The **brightness** of an image is the average intensity of all its pixels

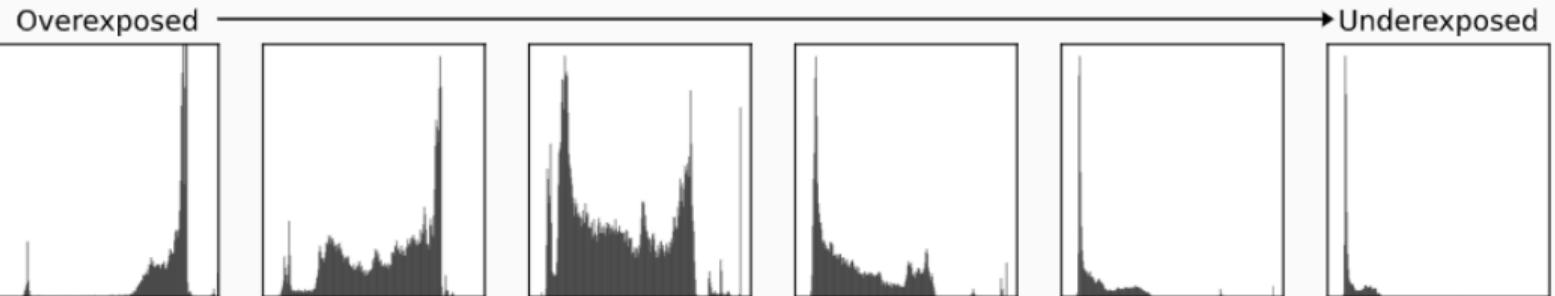
$$B(I) = \frac{1}{w * h} \sum_{r=1}^h \sum_{c=1}^w I(r, c)$$

Where  $w$  and  $h$  are the width and height of the image, and  $I(r, c)$  is the intensity of the pixel at row  $r$  and column  $c$ .

## Brightness example

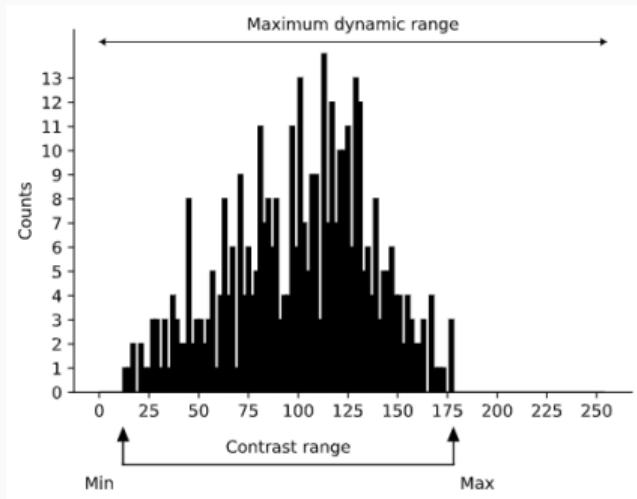


Same X-ray, decreasing brightness - Image from Veldkamp et al., 2009



# Contrast

**Contrast** is the difference in luminance or colour that makes an object distinguishable (Wikipedia). It measures the relationship between light and dark pixels.



# Contrast

**Contrast** is the difference in luminance or colour that makes an object distinguishable (Wikipedia). It measures the relationship between light and dark pixels. Many different definitions in the literature (see Peli, *J. Opt. Soc. Am. A*, 1990)

**Weber contrast** (range  $0 \rightarrow \infty$ )

$$\frac{I_{obj} - I_{bg}}{I_{bg}}$$

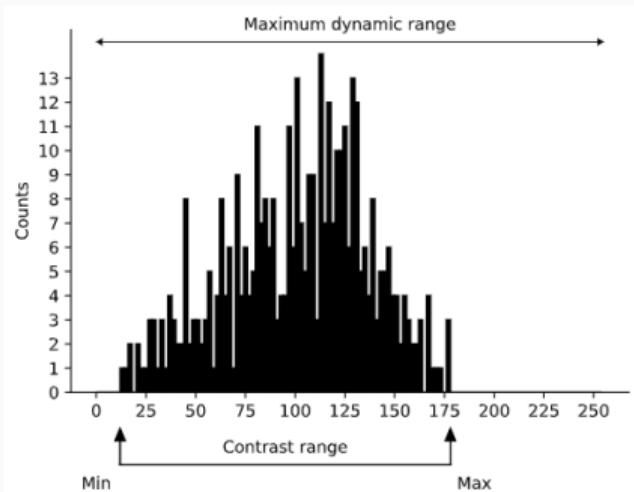
**Michelson contrast** (range  $0 \rightarrow 1$ )

$$\frac{I_{max} - I_{min}}{I_{max} + I_{min}}$$

**RMS contrast**

$$\sqrt{\frac{1}{n-1} \sum_{i=1}^n (x_i - \bar{x})^2}$$

where  $x$  is the normalised intensity  $0 \leq x \leq 1$



## **Point operations**

---

## Point operations and neighborhood operations

Manipulating pixel intensities allows enhancement of images.

Two types of operations:

- Point operations - Change pixel intensity based only on its value
- Neighborhood operations - Change pixel intensity based on the intensity of the pixel and its neighbours.

## Point operations and neighborhood operations

Manipulating pixel intensities allows enhancement of images.

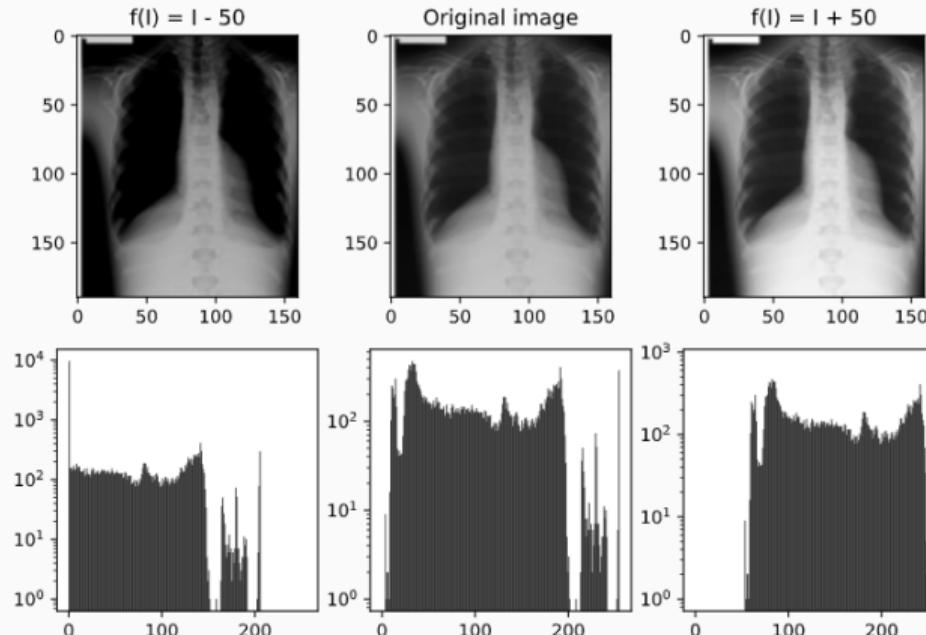
Two types of operations:

- Point operations - Change pixel intensity based only on its value
- Neighborhood operations - Change pixel intensity based on the intensity of the pixel and its neighbours.

So, the resulting intensity  $I'$  of a point operation is  $I'_{(x,y)} = f(I_{x,y})$ .

## Manipulating brightness

We can easily change an image brightness by using a point operation  $f(I_{x,y}) = I_{x,y} + \text{offset}$ , which shifts the image histogram to the left or to the right.



**What is happening to the histograms?**

## Example code for brightness manipulation

```
from skimage import img_as_float, img_as_ubyte

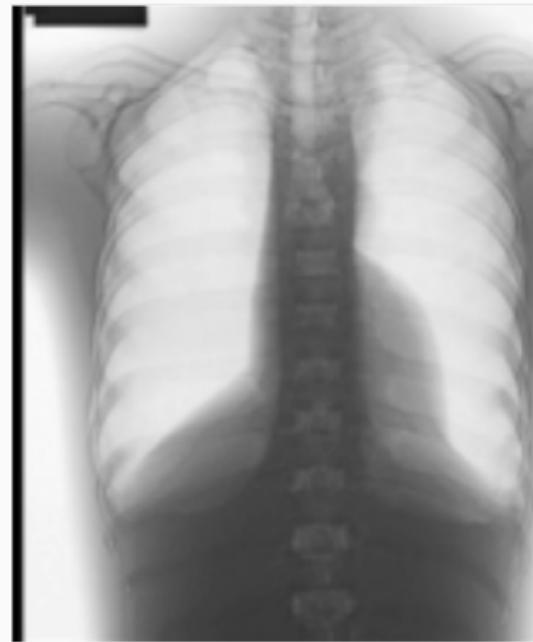
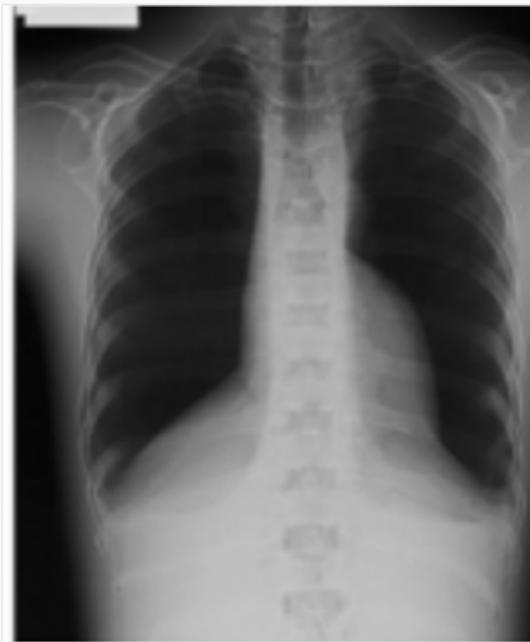
def change_brightness(img, offset):
    """
    Changes the brightness of an image
    img: the image
    offset: the brightness offset to apply
    Returns: the modified image
    """

    img2 = img_as_float(img)
    img2 += offset/255
    img2 = np.clip(img2, 0, 1)
    return (img_as_ubyte(img2))

img_2 = change_brightness(img, 10)
```

## Your turn!

Can you think of a point operation to invert an image? What happens to the histogram?  
Write your own Python function to do this!



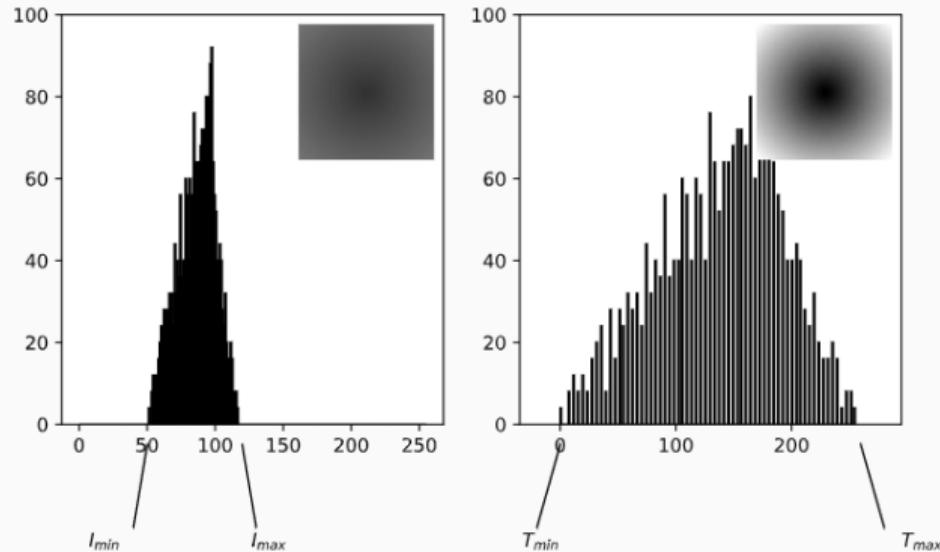
## Manipulating contrast - histogram stretching

A simple way of increasing the contrast of an image is to stretch its histogram.

Simplest form:

$$I' = (I - I_{min}) \frac{T_{max} - T_{min}}{I_{max} - I_{min}} + T_{min}$$

Where  $I_{min}$  and  $I_{max}$  are the minimum and maximum intensity of the image, and  $T_{min}$  and  $T_{max}$  are the minimum and maximum intensity of the target histogram.



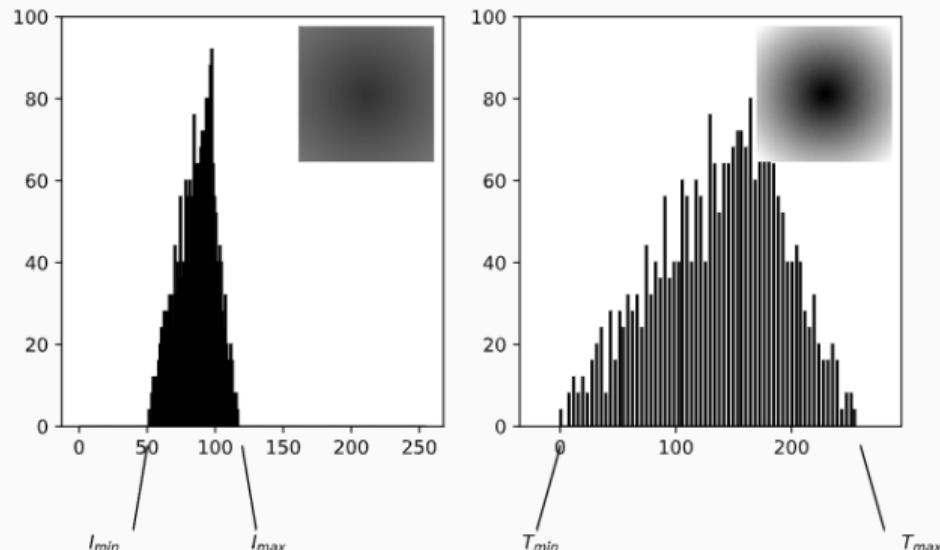
## Manipulating contrast - histogram stretching

A simple way of increasing the contrast of an image is to stretch its histogram.

Simplest form:

$$I' = (I - I_{min}) \frac{T_{max} - T_{min}}{I_{max} - I_{min}} + T_{min}$$

Where  $I_{min}$  and  $I_{max}$  are the minimum and maximum intensity of the image, and  $T_{min}$  and  $T_{max}$  are the minimum and maximum intensity of the target histogram.



This calculation can become meaningless if we have even a single very light or dark pixel, so we can use the 2<sup>nd</sup> and 98<sup>th</sup> percentiles instead of  $I_{min}$  and  $I_{max}$  values.

## Histogram stretching in Python

```
import numpy as np
from skimage.exposure import rescale_intensity

img_rescale = rescale_intensity(img, in_range=(0, 100), out_range=(0, 255))
```

## Histogram stretching in Python

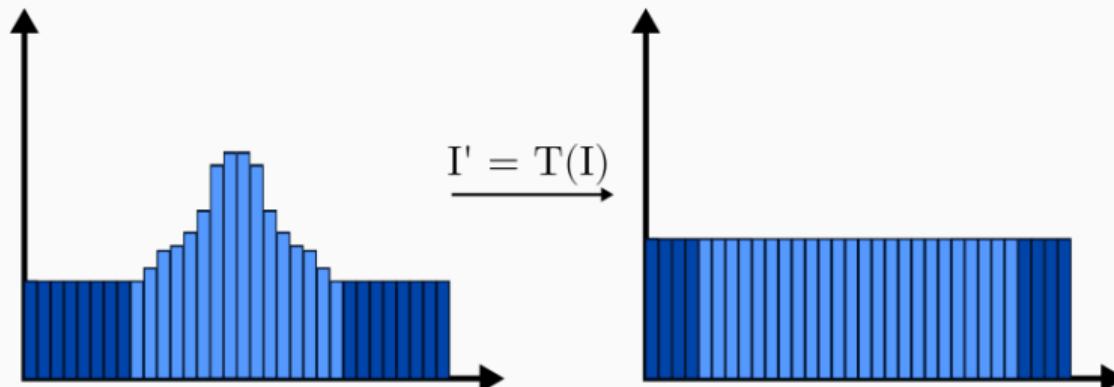
```
import numpy as np
from skimage.exposure import rescale_intensity

img_rescale = rescale_intensity(img, in_range=(0, 100), out_range=(0, 255))
# Or, use the percentile version:
# Calculate the 2nd and 98th percentiles
p2, p98 = np.percentile(img, (2, 98))
img_rescale = rescale_intensity(img, in_range=(p2, p98), out_range=(0, 255))
```

## Histogram equalization

Histogram equalization increases contrast by manipulating pixel intensities.

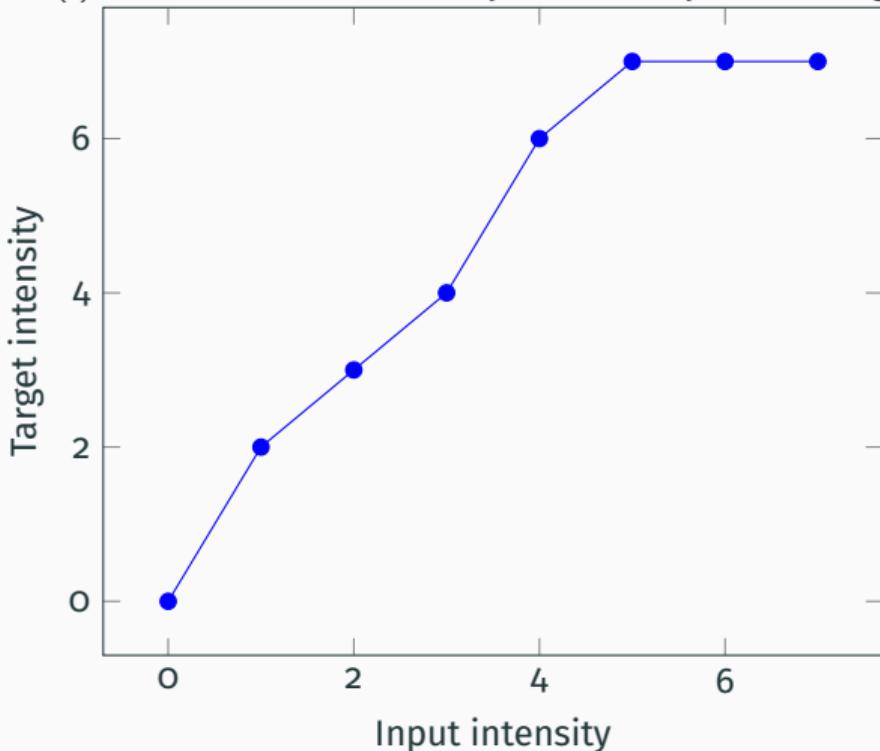
The idea is to generate a **histogram transfer function**  $T(I)$  that will map the image intensities to a target, uniform histogram.



## Choosing $T(I)$

How do we choose  $T(I)$ ?

$T(I)$  should be monotonically (not strictly) increasing, to preserve the intensity rank of the pixels.

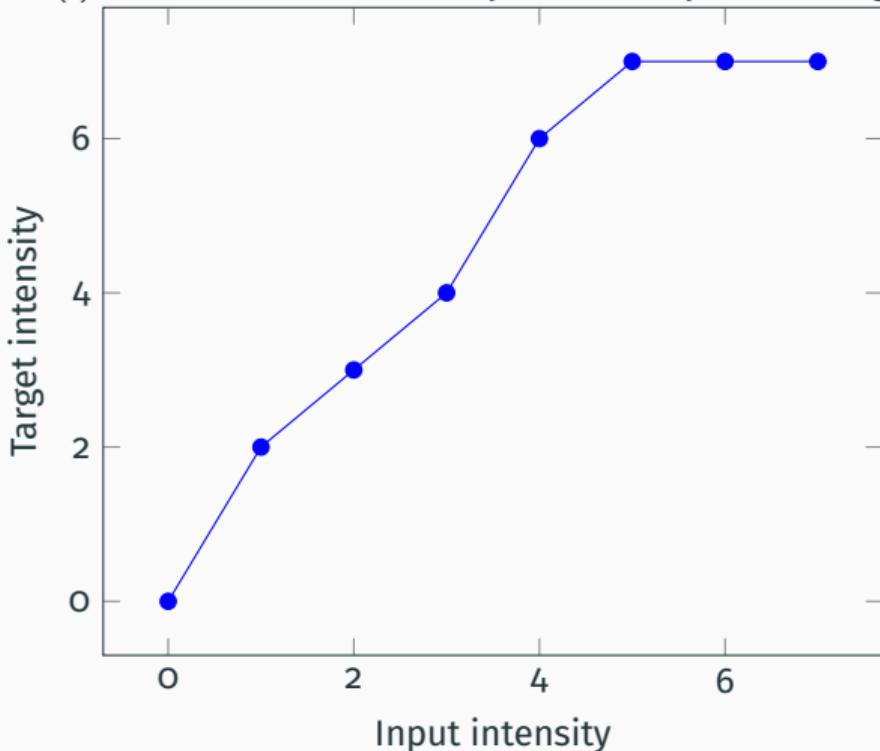


In this example we would change any pixel of the original image with intensity 3 to intensity 4, any with intensity 6 to intensity 7 and so on.

## Choosing $T(I)$

How do we choose  $T(I)$ ?

$T(I)$  should be monotonically (not strictly) increasing, to preserve the intensity rank of the pixels.

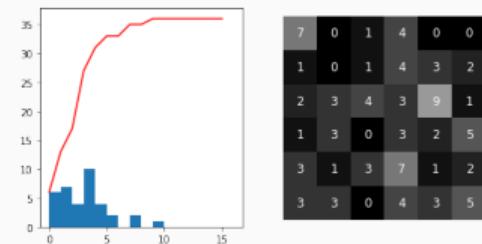


In this example we would change any pixel of the original image with intensity 3 to intensity 4, any with intensity 6 to intensity 7 and so on. It turns out that using the CDF (or the cumulative histogram, for the discrete case) is the solution! Look here if you want a mathematical proof!

## Histogram equalization - an example

Assume you have a 4-bit image (dynamic range 0-15) with the following histogram

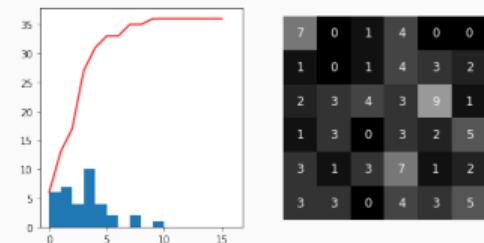
| Pixel value | Number of pixels | Cumulative |
|-------------|------------------|------------|
| 0           | 6                | 6          |
| 1           | 7                | 13         |
| 2           | 4                | 17         |
| 3           | 10               | 27         |
| 4           | 4                | 31         |
| 5           | 2                | 33         |
| 6           | 0                | 33         |
| 7           | 2                | 35         |
| 8           | 0                | 35         |
| 9           | 1                | 36         |
| 10          | 0                | 36         |
| 11          | 0                | 36         |
| 12          | 0                | 36         |
| 13          | 0                | 36         |
| 14          | 0                | 36         |
| 15          | 0                | 36         |



## Histogram equalization - an example

Assume you have a 4-bit image (dynamic range 0-15) with the following histogram

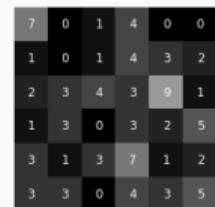
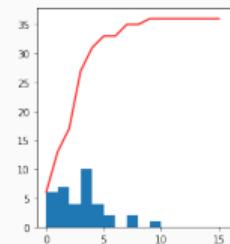
| Pixel value | Number of pixels | Cumulative | Cum. normalized ( $h_n$ ) |
|-------------|------------------|------------|---------------------------|
| 0           | 6                | 6          | 0.17                      |
| 1           | 7                | 13         | 0.36                      |
| 2           | 4                | 17         | 0.47                      |
| 3           | 10               | 27         | 0.75                      |
| 4           | 4                | 31         | 0.86                      |
| 5           | 2                | 33         | 0.92                      |
| 6           | 0                | 33         | 0.92                      |
| 7           | 2                | 35         | 0.97                      |
| 8           | 0                | 35         | 0.97                      |
| 9           | 1                | 36         | 1.00                      |
| 10          | 0                | 36         | 1.00                      |
| 11          | 0                | 36         | 1.00                      |
| 12          | 0                | 36         | 1.00                      |
| 13          | 0                | 36         | 1.00                      |
| 14          | 0                | 36         | 1.00                      |
| 15          | 0                | 36         | 1.00                      |



# Histogram equalization - an example

Assume you have a 4-bit image (dynamic range 0-15) with the following histogram

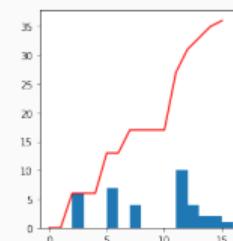
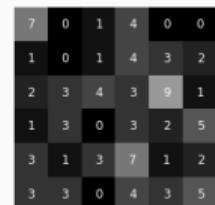
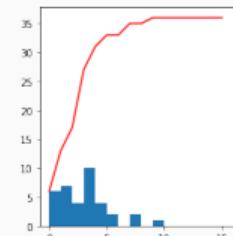
| Pixel value | Number of pixels | Cumulative | Cum. normalized ( $h_n$ ) | New pixel value $\lfloor(h_n * I_{max})\rfloor$ |
|-------------|------------------|------------|---------------------------|---|
| 0           | 6                | 6          | 0.17                      | 2   |
| 1           | 7                | 13         | 0.36                      | 5   |
| 2           | 4                | 17         | 0.47                      | 7   |
| 3           | 10               | 27         | 0.75                      | 11  |
| 4           | 4                | 31         | 0.86                      | 12  |
| 5           | 2                | 33         | 0.92                      | 13  |
| 6           | 0                | 33         | 0.92                      | 13  |
| 7           | 2                | 35         | 0.97                      | 14  |
| 8           | 0                | 35         | 0.97                      | 14  |
| 9           | 1                | 36         | 1.00                      | 15  |
| 10          | 0                | 36         | 1.00                      | 15  |
| 11          | 0                | 36         | 1.00                      | 15  |
| 12          | 0                | 36         | 1.00                      | 15  |
| 13          | 0                | 36         | 1.00                      | 15  |
| 14          | 0                | 36         | 1.00                      | 15  |
| 15          | 0                | 36         | 1.00                      | 15  |



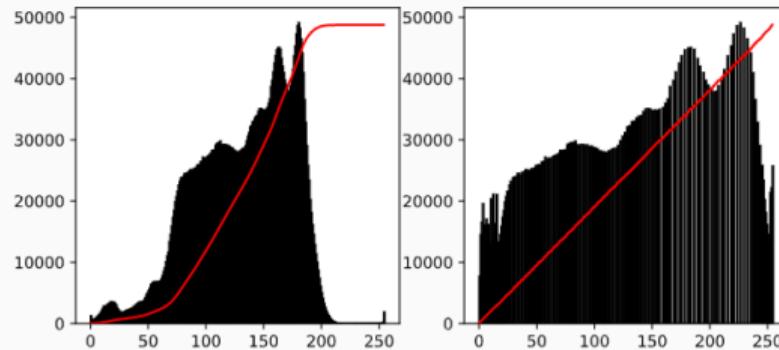
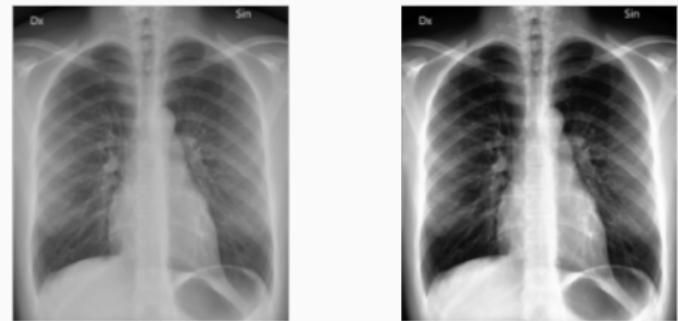
# Histogram equalization - an example

Assume you have a 4-bit image (dynamic range 0-15) with the following histogram

| Pixel value | Number of pixels | Cumulative | Cum. normalized ( $h_n$ ) | New pixel value $\lfloor(h_n * I_{max})\rfloor$ |
|-------------|------------------|------------|---------------------------|---|
| 0           | 6                | 6          | 0.17                      | 2   |
| 1           | 7                | 13         | 0.36                      | 5   |
| 2           | 4                | 17         | 0.47                      | 7   |
| 3           | 10               | 27         | 0.75                      | 11  |
| 4           | 4                | 31         | 0.86                      | 12  |
| 5           | 2                | 33         | 0.92                      | 13  |
| 6           | 0                | 33         | 0.92                      | 13  |
| 7           | 2                | 35         | 0.97                      | 14  |
| 8           | 0                | 35         | 0.97                      | 14  |
| 9           | 1                | 36         | 1.00                      | 15  |
| 10          | 0                | 36         | 1.00                      | 15  |
| 11          | 0                | 36         | 1.00                      | 15  |
| 12          | 0                | 36         | 1.00                      | 15  |
| 13          | 0                | 36         | 1.00                      | 15  |
| 14          | 0                | 36         | 1.00                      | 15  |
| 15          | 0                | 36         | 1.00                      | 15  |



## On a real image...



Note how, in practice, we often do not have a uniform histogram.

## Beyond histogram equalization...

Histogram equalization can result in increased noise and "unrealistic" looking images. Useful e.g. for X-rays, but other strategies exist

## Beyond histogram equalization...

Histogram equalization can result in increased noise and "unrealistic" looking images. Useful e.g. for X-rays, but other strategies exist

**Adaptive histogram equalization** (AHE) computes several histograms of different regions of the image to improve local contrast. Drawback: increased noise in areas with low contrast.

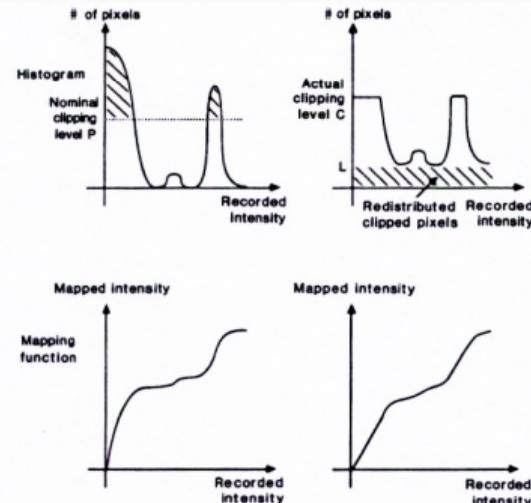
## Beyond histogram equalization...

Histogram equalization can result in increased noise and "unrealistic" looking images. Useful e.g. for X-rays, but other strategies exist

**Adaptive histogram equalization (AHE)** computes several histograms of different regions of the image to improve local contrast. Drawback: increased noise in areas with low contrast.

**Contrast Limited Adaptive Equalization (CLAHE)** solves this by reducing the amount of contrast enhancement. This is done by clipping the histogram to a maximum before calculating the cumulative histogram.

This is implemented in the  
`skimage.exposure.equalize_adapthist` function.



## Histogram equalization - Scikit image

```
from skimage.exposure import equalize_hist, equalize_adapthist  
  
img_equalized = equalize_hist(img)  
# Or, with CLAHE  
img_equalized_CLAHE = equalize_adapthist(img, clip_limit=0.03)
```

Original



equalize\_hist



equalize\_adapthist (clip\_limit 0.02)



equalize\_adapthist (clip\_limit 0.04)



- Histograms are an important tool for inspecting image quality
- Optimal brightness and contrast can help reveal specific image features
- Manipulation of image histograms e.g. through point operations is one important tool for image enhancement