



浙江大学爱丁堡大学联合学院

ZJU-UoE Institute

## Lecture 17 - Hyperparameter tuning

---

Nicola Romanò - [nicola.romano@ed.ac.uk](mailto:nicola.romano@ed.ac.uk)

- Give examples of hyperparameters and explain what are the issues with hyperparameters in deep learning.
- Discuss and compare different methods of tuning hyperparameters.
- Implement these methods in Python.



## Introduction

---

# What are hyperparameters?

**Hyperparameters** are values that are used by a ML model to control the learning process.

Examples include:

- Structural parameters
  - Number of layers
  - Number of units in each layer
  - Number of filters, their kernel size, stride and padding
  - Type of activation function
  - ...
- Learning parameters
  - Type of optimizer
  - Learning rate  $\alpha$
  - Schedule of learning rate
  - Loss function
  - Dropout rate
  - Weight regularization
  - ...
- ...

# What are hyperparameters?

**Hyperparameters** are values that are used by a ML model to control the learning process.

Examples include:

- Structural parameters
  - Number of layers
  - Number of units in each layer
  - Number of filters, their kernel size, stride and padding
  - Type of activation function
  - ...
- Learning parameters
  - Type of optimizer
  - Learning rate  $\alpha$
  - Schedule of learning rate
  - Loss function
  - Dropout rate
  - Weight regularization
  - ...
- ...

Other **parameters** such as weights and biases are learned by the network during training; these will be affected by the choice of hyperparameters.

## Problems with choosing hyperparameters

- Deep network have an extremely large number of hyperparameters.
- The search space is extremely large but often the "good solution" only covers a small region of this space.
- We cannot try all possible combinations.
- No definitive strategy exists.

So... how do we choose hyperparameters?

## Choosing hyperparameters

---

## **Choosing hyperparameters**

---

### **Manual methods**



## What did others do?

We can look at the hyperparameters of published models that performed similar tasks and start from there.

### *Pros*

- If others have used this successfully, there are good chances that might work.
- Computationally fast! :)

## What did others do?

We can look at the hyperparameters of published models that performed similar tasks and start from there.

### *Pros*

- If others have used this successfully, there are good chances that might work.
- Computationally fast! :)

### *Cons*

- It might be difficult to find a model doing exactly the same task.

## Manual tweaking

- Start with a reasonable choice of hyperparameters
- Modify one parameter at a time to improve model performance

- Start with a reasonable choice of hyperparameters
- Modify one parameter at a time to improve model performance

### *Pros*

- It generates high-quality results, when done by an expert.
- Humans can integrate knowledge from several sources and previous experience to make the best changes to hyperparameters.

- Start with a reasonable choice of hyperparameters
- Modify one parameter at a time to improve model performance

### *Pros*

- It generates high-quality results, when done by an expert.
- Humans can integrate knowledge from several sources and previous experience to make the best changes to hyperparameters.

### *Cons*

- It is not obvious to find a good starting point.
- It is extremely time-consuming.
- There is no guarantee to find the optimal model.
- Difficult process to replicate.


## **Choosing hyperparameters**

---

### **Automated methods**

# Grid search

**Grid search** is the simplest method to use and is an exhaustive search over a defined parameter search space.




Grid search - In this example we tune hyperparameters  $x_1$  and  $x_2$  in the range  $[0, 1]$  - source Wikipedia

# Grid search

**Grid search** is the simplest method to use and is an exhaustive search over a defined parameter search space.

## *Pros*

- It is easy to implement.
- Exhaustive search - guaranteed to find the optimal hyperparameters.
- Reproducible.




Grid search - In this example we tune hyperparameters  $x_1$  and  $x_2$  in the range  $[0, 1]$  - source Wikipedia



# Grid search

**Grid search** is the simplest method to use and is an exhaustive search over a defined parameter search space.



Grid search - In this example we tune hyperparameters  $x_1$  and  $x_2$  in the range  $[0, 1]$  - source Wikipedia

## Pros

- It is easy to implement.
- Exhaustive search - guaranteed to find the optimal hyperparameters.
- Reproducible.


## Cons

- It is computationally expensive -> early stop and parallelism help.
- Need to choose a range.
- It does not scale well; 5 hyperparameters with 10 values each, give  $10^5$  combinations.
- A lot of wasted computation time to sample low-accuracy combinations.

We will now see a simple implementation of grid search in Python.  
See the attached `GridSearch.ipynb` notebook.

## Random search


In **random search**, we randomly sample hyperparameters from the search space, rather than exploring a fixed grid.



Random search - In this example we tune hyperparameters  $x_1$  and  $x_2$  in the range  $[0, 1]$ . Random points are chosen in the grid space - source Wikipedia

# Random search

In **random search**, we randomly sample hyperparameters from the search space, rather than exploring a fixed grid.




Random search - In this example we tune hyperparameters  $x_1$  and  $x_2$  in the range  $[0, 1]$ . Random points are chosen in the grid space - source Wikipedia

## Pros

- Still easy to implement.
- Helps when dealing with large search spaces.

# Random search

In **random search**, we randomly sample hyperparameters from the search space, rather than exploring a fixed grid.



Random search - In this example we tune hyperparameters  $x_1$  and  $x_2$  in the range  $[0, 1]$ . Random points are chosen in the grid space - source Wikipedia

## Pros


- Still easy to implement.
- Helps when dealing with large search spaces.

## Cons

- Less reproducible.
- Not guaranteed to hit the optimum.
- Need to choose a range for the hyperparameters.

# Bayesian optimization


**Bayesian** optimization, builds a probability model of the objective function; it uses this model to select the most promising set of hyperparameters. At each iteration, it updates its underlying model, and repeats the process.



Bayesian optimization - In this example we tune hyperparameters  $x_1$  and  $x_2$  in the range  $[0, 1]$ . Bayesian optimization uses a statistical model to choose the next set of hyperparameters *smartly* based on previous results - source Wikipedia

# Bayesian optimization

**Bayesian** optimization, builds a probability model of the objective function; it uses this model to select the most promising set of hyperparameters. At each iteration, it updates its underlying model, and repeats the process.




Bayesian optimization - In this example we tune hyperparameters  $x_1$  and  $x_2$  in the range  $[0, 1]$ . Bayesian optimization uses a statistical model to choose the next set of hyperparameters *smartly* based on previous results - source Wikipedia

## Pros

- Efficient search of hyperparameter space.
- Performs better than grid and random search.

# Bayesian optimization

**Bayesian** optimization, builds a probability model of the objective function; it uses this model to select the most promising set of hyperparameters. At each iteration, it updates its underlying model, and repeats the process.



Bayesian optimization - In this example we tune hyperparameters  $x_1$  and  $x_2$  in the range  $[0, 1]$ . Bayesian optimization uses a statistical model to choose the next set of hyperparameters *smartly* based on previous results - source Wikipedia

## Pros

- Efficient search of hyperparameter space.
- Performs better than grid and random search.

## Cons

- Can be very computationally expensive when using a lot of hyperparameters.
- Complex to implement.



## Early-stopping based methods - Successive Halving

Early-stopping methods work by early stopping of evaluations on combinations of parameters that are not promising.

**Successive Halving** algorithm

- Choose a random set of  $N$  hyperparameters configurations

Early-stopping methods work by early stopping of evaluations on combinations of parameters that are not promising.

**Successive Halving** algorithm

- Choose a random set of  $N$  hyperparameters configurations
- Set a maximum “budget”  $B$  (e.g. number of epochs, time,...) to evaluate these configurations.

Early-stopping methods work by early stopping of evaluations on combinations of parameters that are not promising.

**Successive Halving** algorithm

- Choose a random set of  $N$  hyperparameters configurations
- Set a maximum “budget”  $B$  (e.g. number of epochs, time,...) to evaluate these configurations.
- Each configuration takes up  $B/N$  of the budget.

Early-stopping methods work by early stopping of evaluations on combinations of parameters that are not promising.

**Successive Halving** algorithm

- Choose a random set of  $N$  hyperparameters configurations
- Set a maximum “budget”  $B$  (e.g. number of epochs, time,...) to evaluate these configurations.
- Each configuration takes up  $B/N$  of the budget.
- When the budget is exhausted, discard the worst-performing half of the configurations.

Early-stopping methods work by early stopping of evaluations on combinations of parameters that are not promising.

### **Successive Halving** algorithm

- Choose a random set of  $N$  hyperparameters configurations
- Set a maximum “budget”  $B$  (e.g. number of epochs, time,...) to evaluate these configurations.
- Each configuration takes up  $B/N$  of the budget.
- When the budget is exhausted, discard the worst-performing half of the configurations.
- Continues until only one configuration remains.

Early-stopping methods work by early stopping of evaluations on combinations of parameters that are not promising.

### **Successive Halving** algorithm

- Choose a random set of  $N$  hyperparameters configurations
- Set a maximum “budget”  $B$  (e.g. number of epochs, time,...) to evaluate these configurations.
- Each configuration takes up  $B/N$  of the budget.
- When the budget is exhausted, discard the worst-performing half of the configurations.
- Continues until only one configuration remains.

Early-stopping methods work by early stopping of evaluations on combinations of parameters that are not promising.

**Successive Halving** algorithm

- Choose a random set of  $N$  hyperparameters configurations
- Set a maximum “budget”  $B$  (e.g. number of epochs, time,...) to evaluate these configurations.
- Each configuration takes up  $B/N$  of the budget.
- When the budget is exhausted, discard the worst-performing half of the configurations.
- Continues until only one configuration remains.

**Problem:** do we consider a small  $N$  with a large budget for each configuration, or a large  $N$  with a small budget?

**Hyperband** extends the Successive Halving algorithm by running multiple *brackets* of hyperparameters.

It performs a grid search on the number of configurations  $N$  that we are going to choose.

It can be 10-20x faster than Bayesian optimization.



## KerasTuner

---

The **Keras Tuner** library (install via `pip install keras-tuner`) provides a simple way to tune hyperparameters in a Keras model.

It implements the random search, Bayesian optimization and Hyperband strategies.

In order to use Keras Tuner we need to have a function that generates and returns our Keras model. We will pass that function to the tuner class of our choice.

## KerasTuner example

```
import tensorflow.keras as keras
import keras_tuner as kt

# hp is an object that contains all the hyperparameters
# it will be automatically passed by the tuner
def build_model(hp):
    model = keras.Sequential()
    # Create an integer hyperparameter, going from 10 to 100
    # We can also have hp.Float, hp.Choice, hp.Bool
    n_units = hp.Int('units', min_value=10, max_value=100, step_size=10)
    model.add(keras.layers.Conv2D(n_units, (3, 3), activation='relu',
                                   input_shape=(256, 256, 1)))
    ...
    return model
```

```
rs = kt.tuners.RandomSearch(build_model,
                             objective='val_accuracy',
                             max_trials=10,
                             directory="output_dir")

# Instead of calling model.fit we call tuner.search
rs.search(x_train, y_train,
          x_test, y_test,
          epochs=10, batch_size=32)

best_hyperparameter = rs.get_best_hyperparameters(1)[0]
best_model = rs.get_best_models(1)[0]
```

We can now use a model with the best configuration for our task!